

SUMMARY

A Java/Spring backend developer who turns complex field operations into systems users can run without interruption. I replaced paper-based school inspections with a flow that connects scheduling, checklists, signatures, and re-inspection history, and moved long-running bulk attachment jobs to async processing so users can watch progress and download results. I align with operations staff on the real task order, then design the UI, server, and DB together — owning testing, deployment, and operations after release.

SKILLS

Backend	Java, Spring Boot, Spring MVC, Spring Security, MyBatis
Database	Oracle, MariaDB, MySQL
Frontend	WORK JavaScript, Vue, HTML, CSS, JSP, WebSquare5, jQuery PROJECTS TypeScript, React, Next.js, Redux Toolkit, Pinia
Infra · Delivery	Linux, Tomcat, Jenkins, Gradle, Maven, Git, SVN

EXPERIENCE

TG Narae LLC Web Developer · Solutions Team

Jun 2024 – Present · Full-time

B2B Partner Portal (PPS)

Feb 2025 – Present

A B2B portal where headquarters and about 500 partner companies handle partner registration and contracts, plus training, certification, documentation, and evaluation for 824 field engineers (CEs). I coordinate requirements with operations and business staff, analyze existing data flows, and own DB modeling, HTTP API and SQL design and development, external-system integrations, user acceptance, and Jenkins deployment and operations.

Tech | Java 21 · Spring Boot · MyBatis · MariaDB · Oracle · Vue

PPS Highlights

Moving bulk attachment downloads to async processing

Problem	Collecting and zipping attachments for 300–400 engineers in a single request held the screen for minutes, with no way to tell whether the job was running or had failed.
Decision	The real issue was the request and the job being tied into one flow, not compression speed — so I split the job out and exposed its status.
Implementation	A Spring async job handles compression; the request API returns a job ID and a status API reports running / done / failed. Job state lives in a Hazelcast distributed map, so both servers report the same progress and the job survives a refresh.
Result	Freed HTTP requests from long-running zip jobs and gave users a flow where they track progress and download only finished results.

Unifying the split headquarters account provisioning process

Problem	Accounts had to be created separately in the internal system and the portal, with the portal ID entered directly into the DB — risking omissions and mismatched records.
Decision	Entering data into two systems separately could not prevent omissions, so the portal became the single provisioning gateway.
Implementation	Combined ID creation, organization validation, initial credential issuance, and the internal-system integration into one procedure; if the integration fails, the portal save rolls back in the same transaction.
Result	Collapsed the three-step process — internal account, portal account, manual DB entry — into a single screen, and added inventory and service-ownership checks before account changes to block wrongful deactivations.

PPS Highlights · continued

Rate-limiting password-reset codes across two servers

Problem	Password reset sends a verification code by notification message, but nothing stopped repeated clicks or automated requests — and with two servers, per-server in-memory limits fall apart once requests are load-balanced.
Decision	The limit state had to be shared between servers, and identifying values such as account and phone number must not be stored raw as limiter keys.
Implementation	Stored TTL cooldowns in a Hazelcast distributed map with putIfAbsent-based atomic acquisition, so only one of any concurrent requests sends a code. Limiter keys are SHA-256 hashes of account + phone number, environments without Hazelcast fall back to a local map, and codes come from SecureRandom.
Result	The same limit applies whichever server receives the request, and users see the remaining wait time. Boundary conditions — acquisition right after expiry, concurrent races — are pinned by 30 related unit tests.

Centralizing notification-message policy

Problem	Recipients and send conditions were written into each feature's source, so a policy change meant editing several code paths.
Decision	Separate the frequently changing rules from the shared send logic, so policy changes no longer require a deployment.
Implementation	Recipients and send conditions moved to DB-managed configuration; a shared send service looks up the policy for each message type. Key conditions and exception paths are pinned by unit tests.
Result	Five send points across four business services now go through one shared service; a policy change is a DB configuration update verified by four unit tests.

Education Device Management System (TSMS)

Sep 2025 – Present

Manages about 110,000 education devices for the Seoul Metropolitan Office of Education's one-device-per-student program — registration through delivery, installation, repairs, and inspections, with full history. I align business rules and reference data with operations staff and own DB modeling, HTTP API and SQL design, screen and server development, external integrations, user acceptance, and deployment.

Tech | WebSquare5 · JavaScript · JSP · Java 11 · Spring MVC · MyBatis · MariaDB

TSMS Highlights

Removing exposed API keys and unifying external call paths

Problem	25 screens called map, address, and education-administration APIs directly, exposing keys in the browser and requiring per-screen edits whenever integration details changed.
Decision	Rotating the keys would not stop the exposure or the repeated edits from recurring — the external-call responsibility had to move from the screens to the server.
Implementation	Consolidated external calls into a shared server-side API layer, with keys and endpoint URLs held in server configuration. The affected screens were switched to one common path and response shape.
Result	Eliminated browser key exposure across 25 screens and reduced the scattered call and configuration change points to a single server path.

Building secondhand-marketplace monitoring

Problem	The operations team searched multiple marketplaces by keyword and region by hand and tracked suspicious listings manually. Security policy and site restrictions ruled out direct search or crawling by the system.
Decision	With direct calls to external sites off the table, keep human review and let the server analyze only the pasted listing URL string.
Implementation	The server identifies the site and listing ID from the URL to detect duplicates, with per-site rules managed in reference data. Links, captures, sale details, and handling history are connected on one screen.
Result	Systematized the recording, duplicate checks, and case history that follow a search — without the system ever contacting external sites directly.

TSMS Highlights · continued

Bulk device registration checks and QR issuance

Problem	Bulk spreadsheet uploads could include serial numbers missing from production records, or devices that were already registered.
Decision	Move validation ahead of the save — checking reference data and duplicates before registration instead of cleaning up afterward.
Implementation	Serial numbers are checked against production records, and already-registered devices return their current school and user. QR codes carry an encrypted token instead of internal identifiers.
Result	Applied to the QR issuance flow for 108,237 of 111,593 devices, blocking wrong registrations and identifier exposure up front. (as of Jul 2026)

Field inspection and re-inspection history

Problem	Paper-based school inspections made overall status hard to see, and the 2nd and 3rd revisit history for lost or missing devices was not connected.
Decision	School-level completion flags could not carry revisits forward — inspections needed per-device rounds linked to results.
Implementation	Unfinished devices roll into the next round; required fields and e-signatures are validated before confirmation documents and result files are issued.
Result	Rolled out to 119 target schools; 21 schools completed, with 3,800 device inspections recorded in the system. (as of Jul 2026)

PERSONAL PROJECTS

ReachRich · personal investment research & operations platform Mar 2026 – Present · Personal

Rebuilt a personal investment platform into six modules — account tracking on the Toss Securities Open API, KRX data collection, strategy validation, paper trading, and a dashboard. Implemented date-keyed idempotent persistence with partial-failure isolation, FastAPI query APIs, a React UI, and GitHub Actions health checks with Telegram alerts; 205 backend tests cover the data collection, validation, and operation logic.

Tech | Python · FastAPI · SQLAlchemy · SQLite · Parquet · React · TypeScript · GitHub Actions

SSAFAST · API spec & testing collaboration tool Apr 2023 – May 2023 · Team of 6

As the frontend / UI-UX developer on a team of six, I built the API spec authoring screens with repeated and nested inputs, and the test result screens showing request success, response contents, and response-time distribution.

Tech | TypeScript · React · Next.js · Redux Toolkit · React Hook Form · TanStack Query

EDUCATION

Training	Samsung SW Academy For Youth (SSAFY), 8th cohort · Web development program	Jul 2022 – Jun 2023
Education	Ajou University · B.S. in e-Business (transfer)	Mar 2018 – Aug 2020
Education	California State University, Chico · Business Administration, attended before transfer	Jan 2014 – May 2015
Certification	SQLD (SQL Developer) · Korea Data Agency	Sep 2024